

BraindumpStudy



Latest updated materials, Daily Updates!



<http://www.braindumpstudy.com>

BraindumpStudy Exam Dumps, High Pass Rate!

Exam : **AWS-Solutions-Architect-Professional**

Title : AWS Certified Solutions Architect - Professional

Vendor : Amazon

Version : DEMO

NO.1 A user is configuring MySQL RDS with PIOPS. What should be the minimum size of DB storage provided by the user?

- A. 1 TB
- B. 50 GB
- C. 5 GB
- D. 100 GB

Answer: D

Explanation:

If the user is trying to enable PIOPS with MySQL RDS, the minimum size of storage should be 100 GB.

Reference: http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_PIOPS.html

NO.2 Select the correct statement about Amazon ElastiCache.

- A. It makes it easy to set up, manage, and scale a distributed in-memory cache environment in the cloud.
- B. It allows you to quickly deploy your cache environment only if you install software.
- C. It does not integrate with other Amazon Web Services.
- D. It cannot run in the Amazon Virtual Private Cloud (Amazon VPC) environment.

Answer: A

Explanation:

ElastiCache is a web service that makes it easy to set up, manage, and scale a distributed in-memory cache environment in the cloud. It provides a high-performance, scalable, and cost-effective caching solution, while removing the complexity associated with deploying and managing a distributed cache environment. With ElastiCache, you can quickly deploy your cache environment, without having to provision hardware or install software.

Reference: <http://docs.aws.amazon.com/AmazonElastiCache/latest/UserGuide/WhatIs.html>

NO.3 How can a user list the IAM Role configured as a part of the launch config?

- A. `as-describe-launch-configs --iam-profile`
- B. `as-describe-launch-configs --show-long`
- C. `as-describe-launch-configs -iam-role`
- D. `as-describe-launch-configs -role`

Answer: B

Explanation:

`As-describe-launch-configs` describes all the launch config parameters created by the AWS account in the

specified region. Generally it returns values, such as Launch Config name, Instance Type and AMI ID.

If

the user wants additional parameters, such as the IAM Profile used in the config, he has to run command:

`as-describe-launch-configs --show-long`

NO.4 What RAID method is used on the Cloud Block Storage back-end to implement a very high level of

reliability and performance?

- A. RAID 1 (Mirror)
- B. RAID 5 (Blocks striped, distributed parity)

C. RAID 10 (Blocks mirrored and striped)

D. RAID 2 (Bit level striping)

Answer: C

Explanation:

Cloud Block Storage back-end storage volumes employs the RAID 10 method to provide a very high level

of reliability and performance.

Reference: http://www.rackspace.com/knowledge_center/product-faq/cloud-block-storage

NO.5 Which of the following is the Amazon Resource Name (ARN) condition operator that can be used within an Identity and Access Management (IAM) policy to check the case-insensitive matching of the ARN?

A. ArnCheck

B. ArnMatch

C. ArnCase

D. ArnLike

Answer: D

Explanation:

Amazon Resource Name (ARN) condition operators let you construct Condition elements that restrict access based on comparing a key to an ARN. ArnLike, for instance, is a case-insensitive matching of the

ARN. Each of the six colon-delimited components of the ARN is checked separately and each can include

a multi-character match wildcard (*) or a single-character match wildcard (?).

Reference:

http://docs.aws.amazon.com/IAM/latest/UserGuide/AccessPolicyLanguage_ElementDescriptions.html

NO.6 You are looking to migrate your Development (Dev) and Test environments to AWS. You have decided to

use separate AWS accounts to host each environment. You plan to link each accounts bill to a Master AWS account using Consolidated Billing. To make sure you Keep within budget you would like to implement a way for administrators in the Master account to have access to stop, delete and/or terminate

resources in both the Dev and Test accounts. Identify which option will allow you to achieve this goal.

A. Create IAM users in the Master account with full Admin permissions. Create cross-account roles in the

Dev and Test accounts that grant the Master account access to the resources in the account by inheriting

permissions from the Master account.

B. Create IAM users and a cross-account role in the Master account that grants full Admin permissions to

the Dev and Test accounts.

C. Create IAM users in the Master account Create cross-account roles in the Dev and Test accounts that

have full Admin permissions and grant the Master account access.

D. Link the accounts using Consolidated Billing. This will give IAM users in the Master account access

to
resources in the Dev and Test accounts

Answer: C

NO.7 Your company is getting ready to do a major public announcement of a social media site on AWS. The website is running on EC2 instances deployed across multiple Availability Zones with a Multi-AZ RDS MySQL Extra Large DB Instance. The site performs a high number of small reads and writes per second and relies on an eventual consistency model. After comprehensive tests you discover that there is contention on RDS MySQL. Which are the best approaches to meet these requirements? (Choose 2 answers)

- A. Deploy ElastiCache in-memory cache running in each availability zone
- B. Implement sharding to distribute load to multiple RDS MySQL instances
- C. Increase the RDS MySQL Instance size and Implement provisioned IOPS
- D. Add an RDS MySQL read replica in each availability zone

Answer: A, C

NO.8 The user has provisioned the PIOPS volume with an EBS optimized instance. Generally speaking, in which I/O chunk should the bandwidth experienced by the user be measured by AWS?

- A. 128 KB
- B. 256 KB
- C. 64 KB
- D. 32 KB

Answer: B

Explanation:

IOPS are input/output operations per second. Amazon EBS measures each I/O operation per second (that is 256 KB or smaller) as one IOPS.

Reference: <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-io-characteristics.html>

NO.9 What is the role of the PollForTask action when it is called by a task runner in AWS Data Pipeline?

- A. It is used to retrieve the pipeline definition.
- B. It is used to report the progress of the task runner to AWS Data Pipeline.
- C. It is used to receive a task to perform from AWS Data Pipeline.
- D. It is used to inform AWS Data Pipeline of the outcome when the task runner completes a task.

Answer: C

Explanation:

Task runners call PollForTask to receive a task to perform from AWS Data Pipeline. If tasks are ready in

the work queue, PollForTask returns a response immediately. If no tasks are available in the queue, PollForTask uses long-polling and holds on to a poll connection for up to 90 seconds, during which time

any newly scheduled tasks are handed to the task agent. Your remote worker should not call PollForTask

again on the same worker group until it receives a response, and this may take up to 90 seconds.
Reference: http://docs.aws.amazon.com/datapipeline/latest/APIReference/API_PollForTask.html

NO.10 An organization is setting up a highly scalable application using Elastic Beanstalk. They are using Elastic Load Balancing (ELB) as well as a Virtual Private Cloud (VPC) with public and private subnets. They have the following requirements:

- . All the EC2 instances should have a private IP
- . All the EC2 instances should receive data via the ELB's.

Which of these will not be needed in this setup?

- A. Launch the EC2 instances with only the public subnet.
- B. Create routing rules which will route all inbound traffic from ELB to the EC2 instances.
- C. Configure ELB and NAT as a part of the public subnet only.
- D. Create routing rules which will route all outbound traffic from the EC2 instances through NAT.

Answer: A

Explanation:

The Amazon Virtual Private Cloud (Amazon VPC) allows the user to define a virtual networking environment in a private, isolated section of the Amazon Web Services (AWS) cloud. The user has complete control over the virtual networking environment. If the organization wants the Amazon EC2 instances to have a private IP address, he should create a public and private subnet for VPC in each Availability Zone (this is an AWS Elastic Beanstalk requirement). The organization should add their public

resources, such as ELB and NAT to the public subnet, and AWS Elastic Beanstalk will assign them unique elastic IP addresses (a static, public IP address). The organization should launch Amazon EC2 instances in a private subnet so that AWS Elastic Beanstalk assigns them non-routable private IP addresses. Now the organization should configure route tables with the following rules:

- . route all inbound traffic from ELB to EC2 instances
- . route all outbound traffic from EC2 instances through NAT

Reference: <http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/AWSHowTo-vpc.html>

NO.11 Which statement is NOT true about a stack which has been created in a Virtual Private Cloud (VPC) in AWS OpsWorks?

- A. Subnets whose instances cannot communicate with the Internet are referred to as public subnets.
- B. Subnets whose instances can communicate only with other instances in the VPC and cannot communicate directly with the Internet are referred to as private subnets.
- C. All instances in the stack should have access to any package repositories that your operating system depends on, such as the Amazon Linux or Ubuntu Linux repositories.
- D. Your app and custom cookbook repositories should be accessible for all instances in the stack.

Answer: A

Explanation:

In AWS OpsWorks, you can control user access to a stack's instances by creating it in a virtual private cloud (VPC). For example, you might not want users to have direct access to your stack's app servers or databases and instead require that all public traffic be channeled through an Elastic Load Balancer.

A VPC consists of one or more subnets, each of which contains one or more instances. Each subnet has

an associated routing table that directs outbound traffic based on its destination IP address.

Instances within a VPC can generally communicate with each other, regardless of their subnet.

Subnets whose instances can communicate with the Internet are referred to as public subnets.

Subnets whose instances can communicate only with other instances in the VPC and cannot communicate directly with the Internet are referred to as private subnets.

AWS OpsWorks requires the VPC to be configured so that every instance in the stack, including instances in private subnets, has access to the following endpoints:

The AWS OpsWorks service, <https://opsworks-instance-service.us-east-1.amazonaws.com>.

Amazon S3

The package repositories for Amazon Linux or Ubuntu 12.04 LTS, depending on which operating system

you specify.

Your app and custom cookbook repositories.

Reference:

<http://docs.aws.amazon.com/opsworks/latest/userguide/workingstacks-vpc.html#workingstacks-vpc-basics>

NO.12 What is a possible reason you would need to edit claims issued in a SAML token?

A. The NamelIdentifier claim cannot be the same as the username stored in AD.

B. Authentication fails consistently.

C. The NamelIdentifier claim cannot be the same as the claim URI.

D. The NamelIdentifier claim must be the same as the username stored in AD.

Answer: A

Explanation:

The two reasons you would need to edit claims issued in a SAML token are:

The NamelIdentifier claim cannot be the same as the username stored in AD, and

The app requires a different set of claim URIs.

Reference:

<https://azure.microsoft.com/en-us/documentation/articles/active-directory-saml-claims-customization/>

NO.13 Regarding Amazon SNS, you can send notification messages to mobile devices through any of the

following supported push notification services, EXCEPT:

A. Microsoft Windows Mobile Messaging (MWMM)

B. Google Cloud Messaging for Android (GCM)

C. Amazon Device Messaging (ADM)

D. Apple Push Notification Service (APNS)

Answer: A

Explanation:

In Amazon SNS, you have the ability to send notification messages directly to apps on mobile devices.

Notification messages sent to a mobile endpoint can appear in the mobile app as message alerts, badge

updates, or even sound alerts. Microsoft Windows Mobile Messaging (MWMM) doesn't exist and is not

supported by Amazon SNS.

Reference: <http://docs.aws.amazon.com/sns/latest/dg/SNSMobilePush.html>

NO.14 Which of the following cannot be done using AWS Data Pipeline?

- A. Regularly access your data where it's stored, transform and process it at scale, and efficiently transfer the results to another AWS service.
- B. Move data between different AWS compute and storage services as well as on-premise data sources at specified intervals.
- C. Generate reports over data that has been stored.
- D. Create complex data processing workloads that are fault tolerant, repeatable, and highly available.

Answer: C

NO.15 You are designing Internet connectivity for your VPC. The Web servers must be available on the Internet.

The application must have a highly available architecture.

Which alternatives should you consider? (Choose 2 answers)

- A. Configure a NAT instance in your VPC Create a default route via the NAT instance and associate itwith all subnets Configure a DNS A record that points to the NAT instance public IP address.
- B. Configure a CloudFront distribution and configure the origin to point to the private IP addresses of your Web servers Configure a Route53 CNAME record to your CloudFront distribution.
- C. Place all your web servers behind ELB Configure a Route53 CNMIE to point to the ELB DNS name.
- D. Assign EIPs to all web servers. Configure a Route53 record set with all EIPs, with health checks and DNS failover.
- E. Configure ELB with an EIP Place all your Web servers behind ELB Configure a Route53 A record that points to the EIP.

Answer: C, D

NO.16 A government client needs you to set up secure cryptographic key storage for some of their extremely

confidential data. You decide that the AWS CloudHSM is the best service for this. However, there seem to

be a few pre-requisites before this can happen, one of those being a security group that has certain ports

open. Which of the following is correct in regards to those security groups?

- A. A security group that has no ports open to your network.
- B. A security group that has only port 3389 (for RDP) open to your network.
- C. A security group that has only port 22 (for SSH) open to your network.
- D. A security group that has port 22 (for SSH) or port 3389 (for RDP) open to your network.

Answer: D

Explanation:

AWS CloudHSM provides secure cryptographic key storage to customers by making hardware securit
y

modules (HSMs) available in the AWS cloud.

AWS CloudHSM requires the following environment before an HSM appliance can be provisioned.

A virtual private cloud (VPC) in the region where you want the AWS CloudHSM service.

One private subnet (a subnet with no Internet gateway) in the VPC. The HSM appliance is provisioned into this subnet.

One public subnet (a subnet with an Internet gateway attached). The control instances are attached to this subnet.

An AWS Identity and Access Management (IAM) role that delegates access to your AWS resources to AWS CloudHSM.

An EC2 instance, in the same VPC as the HSM appliance, that has the SafeNet client software installed.

This instance is referred to as the control instance and is used to connect to and manage the HSM appliance.

A security group that has port 22 (for SSH) or port 3389 (for RDP) open to your network. This security group is attached to your control instances so you can access them remotely.

NO.17 Which of the following statements is correct about the number of security groups and rules applicable for

an EC2-Classic instance and an EC2-VPC network interface?

A. In EC2-Classic, you can associate an instance with up to 5 security groups and add up to 50 rules to a

security group. In EC2-VPC, you can associate a network interface with up to 500 security groups and add up to 100 rules to a security group.

B. In EC2-Classic, you can associate an instance with up to 500 security groups and add up to 50 rules to

a security group. In EC2-VPC, you can associate a network interface with up to 5 security groups and add

up to 100 rules to a security group.

C. In EC2-Classic, you can associate an instance with up to 5 security groups and add up to 100 rules to a

security group. In EC2-VPC, you can associate a network interface with up to 500 security groups and add up to 50 rules to a security group.

D. In EC2-Classic, you can associate an instance with up to 500 security groups and add up to 100 rules

to a security group. In EC2-VPC, you can associate a network interface with up to 5 security groups and

add up to 50 rules to a security group.

Answer: D

Explanation:

A security group acts as a virtual firewall that controls the traffic for one or more instances. When you

launch an instance, you associate one or more security groups with the instance. You add rules to each

security group that allow traffic to or from its associated instances. If you're using EC2-Classic, you must

use security groups created specifically for EC2-Classic. In EC2-Classic, you can associate an instance with up to 500 security groups and add up to 100 rules to a security group. If you're using EC2-VPC, you

must use security groups created specifically for your VPC. In EC2-VPC, you can associate a network interface with up to 5 security groups and add up to 50 rules to a security group.

Reference: <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html>

NO.18 Your system recently experienced down time during the troubleshooting process. You found that a new

administrator mistakenly terminated several production EC2 instances.

Which of the following strategies will help prevent a similar situation in the future?

The administrator still must be able to:

launch, start stop, and terminate development resources.

launch and start production instances.

A. Create an IAM user, which is not allowed to terminate instances by leveraging production EC2 termination protection.

B. Leverage resource based tagging, along with an IAM user which can prevent specific users from terminating production, EC2 resources.

C. Leverage EC2 termination protection and multi-factor authentication, which together require users to

authenticate before terminating EC2 instances

D. Create an IAM user and apply an IAM role which prevents users from terminating production EC2 instances.

Answer: B

NO.19 You are running a successful multitier web application on AWS and your marketing department has asked

you to add a reporting tier to the application. The reporting tier will aggregate and publish status reports

every 30 minutes from user-generated information that is being stored in your web application s database.

You are currently running a Multi-AZ RDS MySQL instance for the database tier. You also have implemented ElastiCache as a database caching layer between the application tier and database tier. Please select the answer that will allow you to successfully implement the reporting tier with as little impact as possible to your database.

A. Continually send transaction logs from your master database to an S3 bucket and generate the reports

off the S3 bucket using S3 byte range requests.

B. Generate the reports by querying the synchronously replicated standby RDS MySQL instance maintained through Multi-AZ.

C. Launch a RDS Read Replica connected to your Multi AZ master database and generate reports by querying the Read Replica.

D. Generate the reports by querying the ElastiCache database caching tier.

Answer: C

NO.20 Your company plans to host a large donation website on Amazon Web Services (AWS). You anticipate a

large and undetermined amount of traffic that will create many database writes. To be certain that you do

not drop any writes to a database hosted on AWS. Which service should you use?

- A. Amazon RDS with provisioned IOPS up to the anticipated peak write throughput.
- B. Amazon Simple Queue Service (SQS) for capturing the writes and draining the queue to write to the database.
- C. Amazon ElastiCache to store the writes until the writes are committed to the database.
- D. Amazon DynamoDB with provisioned write throughput up to the anticipated peak write throughput.

Answer: B